

Chess Position Recognition

Square-level board understanding from schematic chess images

Dataset	Kaggle Chess Positions
Model	Scratch residual CNN in PyTorch
Primary metric	Square accuracy over 64 board cells

Project Goal

- Input: one aligned 400 by 400 RGB chessboard image.
- Output: one label for every board square.
- Label set: empty, white pieces, and black pieces.
- Success is measured primarily by square-level accuracy.

Prediction Target

64

squares per image

Classes

13

empty + 12 pieces

Main Metric Accuracy

square-level correctness

Dataset


- Source: [Chess Positions](#) dataset.
- 100,000 generated schematic board images.
- 80,000 training images and 20,000 test images.
- Each image contains 5–15 pieces, including both kings.
- Images span many board and piece style combinations.
- Filename stems encode labels as FEN ranks with dashes instead of slashes.

Data Contract

Field	Requirement
Location	data/train/ and data/test/
Image	RGB board image, read from disk every run
Label	Filename stem with 8 dash-separated FEN ranks
Split	Random validation split from data/train/; final test uses data/test/
Cache policy	No processed dataset cache is written


Label Parsing

- A filename stem is split into 8 ranks.
- Digits expand to repeated empty squares.
- Piece letters map directly to class ids.
- The result is an 8 x 8 label grid.




qR6-N7- . . . becomes square labels.

Pipeline



Model Architecture

- Implemented directly in PyTorch as BoardCNN.
- Residual convolution blocks learn piece and board features.
- Spatial downsampling aligns features to an 8 by 8 grid.
- A 1 by 1 convolution emits class logits for each square.



Training Setup

Component	Setting
Optimizer	AdamW with weight decay
Loss	Cross-entropy over square labels
Empty handling	Lower empty-class weight to focus on occupied squares
Regularization	Dropout, label smoothing, and label-preserving augmentation
Hardware target	Local RTX 4080 with CUDA mixed precision
Checkpoint rule	Best model selected by validation square accuracy


Metrics

- **Square accuracy:** primary metric over every board square.
- **Occupied-square accuracy:** piece recognition without empty squares.
- **Empty-square accuracy:** empty-cell recognition quality.
- **Board accuracy:** strict metric requiring all 64 squares correct.
- **Per-class precision/recall/F1:** piece-specific behavior.



Generated Outputs

Path	Purpose
output/models/	PyTorch checkpoints and final model
output/predictions/	Square-level predictions in CSV and Parquet
output/reports/	Metrics, class reports, and training history
output/figures/	Training curves, confusion matrix, confidence plots

Training Curves



Evaluation Figures



Results Placeholder

Split	Square Accuracy	Occupied Accuracy	Board Accuracy
Validation	From <code>metrics.csv</code>	After training	After training
Test	From <code>metrics.csv</code>	After training	After training

Rebuild this deck after the long run with just `presentation-build`.

Risks and Mitigations

Risk	Mitigation
Empty-square dominance	Down-weight empty class and track occupied-square accuracy
Style variation	Train on all board and piece styles with color augmentation
Overfitting	Validation split, dropout, label smoothing, early stopping
Slow full training	Mixed precision and RTX 4080 local execution
Report drift	Presentation references generated outputs directly

Reproducibility

- Install environment: `just sync`
- Smoke test: `just smoke`
- Full training: `just run`
- Build docs: `just docs-build`
- Build presentation: `just presentation-build`

The project reads raw images on each run and writes all report artifacts under `output/`.

Conclusion

- The project has been migrated to chess board position recognition.
- The model predicts all 64 square labels with one scratch CNN.
- Accuracy is measured at square, occupied-square, empty-square, and full-board levels.
- The remaining project work is the long RTX 4080 training run and final result interpretation.

Questions

Chess Position Recognition